

Supplemental Note B

Supplementary material for

ECDSA.Fail: Open Autoresearch for Optimizing Elliptic-Curve Point Addition in Shor's Algorithm

The paper cites this document as Supplemental Note B. Section, figure, and table numbers continue those of the main paper, so a reference such as Section 5.3.1 or Table 2 carries the same number here as it does there. References to material outside this note resolve in the main paper or in a sibling Supplemental Note.

B Anatomy of the Agent-Identified Quantum Circuit Optimizations

Section 5.4.4 classified the 400 scored source commits of the public competition into the nine optimization families and two residual buckets of Figure 10. This note supplies the mechanism-level anatomy behind that census: for each family, what the transformation does to the reversible circuit, which correctness obligation licenses it, which factor of the score $S = QT$ (Section 3.2) it moves, and which audited commits attest it. Where a family's marquee instance is already developed in the main text—the Jump-2 Euclidean recurrence (Section 5.3.1), the base-5 transcript codec (Section 5.3.2), specialized Karatsuba squaring (Section 5.3.4), register-shared Euclidean inversion (Section 5.3.5), constant propagation (Section 5.3.6), and dead-code elimination (Section 5.3.7)—we do not repeat that treatment, and record only the technique's placement in the taxonomy, its neighbouring mechanisms, and its commit-level record. The complete per-commit ledger, with each commit's joined leaderboard score, frontier delta, contestant, and date, is the census table of Supplemental Note C.

The archive records landed source transformations more reliably than it records who first conceived them. Accordingly, *agent-identified* denotes optimization episodes surfaced and developed in an agent-mediated workflow; it does not assign exclusive agency, authorship, or novelty to a model. Human direction, community exchange, and imported research remain part of the provenance, and are separated from local implementation claims below.

Why these nine families. The organizing idea is that the scored artifact is, in effect, a compiled straight-line reversible arithmetic program: a fixed schedule of modular additions, multiplications, and inversions over $\mathbb{F}_p$ realized as Toffoli-and-Clifford gates on a bounded qubit register. Read this way, the transformations in the archive are recognizably related to moves a classical optimizing compiler makes on such a program, and six families accordingly carry familiar compiler labels: algorithmic substitution, strength reduction, fusion, windowing (loop blocking with precomputation), constant propagation, and dead-code elimination. The remaining three have familiar classical analogues but are separated

here because reversibility changes their objective or proof duty: an *uncompute strategy* family, because scratch must be returned to $|0\rangle$ without residual phase; a *register and scratch allocation* family, because the location and lifetime of quantum data directly set peak logical width; and a *codec/encoding* family, because reversible decoding and cleanup determine the time–space exchange when resident quantum state is replaced by encoded metadata. The claim is thus not that classical compilers lack allocation or encoding, but that these transformations acquire distinct semantic and resource consequences in a reversible circuit.

This transformation axis is deliberately orthogonal to the two the main text already used: the metric axis—which score factor a change primarily moves—and the primitive axis—which building block it touched (adder, comparator, modular multiplier, Kaliski inverse). Because each commit is assigned the single family naming its *dominant* transformation, the families are a single-label characterization rather than a claim that the underlying mechanisms are disjoint; a submission routinely bundles several—a width truncation co-tuned with a nonce reroll, or a carry-fold gadget that at once removes gates, substitutes a cheaper equivalent, and merges two adjacent operations—and the assignment records only which of them is primary. This boundary was the principal source of disagreement during a second audit, which prompted reassignment of multi-mechanism entries; because the released census preserves only the reconciled primary labels, not the paired pre-reconciliation labels, we do not report an inter-rater reliability statistic. The audit left the qualitative conclusions below unchanged. The taxonomy is likewise open rather than closed: later eras of the search—still wider windowing and, most concretely, a prospective *representation / coordinate change* family (coset, projective, or Kim wide- r representations, characterized in pre-competition research but with no landed competition instance)—may force additional families. We order the subsections below by structural role rather than by frequency, leading with the skeleton-defining windowing family.

Which commits each claim rests on. This note draws on both halves of the source history defined in Section 5.4.4. The census of Figure 10 counts only the scored commits of the public competition and is reconciled against the public ledger in Supplemental Note C; the technique-level discussion inside each family instead draws mainly on the earlier pre-competition commits, where the mechanism first becomes source-auditable.¹

¹ The per-commit Toffoli and peak-qubit figures quoted in the family subsections below are source-reported local counts, typically measured on that technique’s own primitive or working tree at the pre-competition commit where it landed (as recorded in the corresponding source audit), not the global operating points of the headline table in Supplemental Note A. The working-tree peaks these deltas sit on—often in the low thousands of logical qubits—therefore differ from the promoted submissions’ headline widths, and most identifiers are pre-competition development hashes; where a competition submission hash is quoted (e.g., 9a81017), the figures are that submission’s own operating point.

Distinguishing gate removal from re-sampling. One classification boundary cannot be settled from the diff alone and requires running the evaluator. Because the score averages *executed* Toffoli over a Fiat–Shamir sample that the submission’s nonce selects (Section 3.3), and because the low-width routes contain measurement-conditioned Toffolis whose execution is data-dependent, a submission can lower its score in two ways that look identical on paper: by *removing gates*, or by *re-sampling* to a luckier “island” on which the same gates execute fewer times. We separate them by rebuilding each ambiguous submission and comparing its *static* Toffoli-gate count—after nonce-tail and explicit padding changes are excluded—to its parent’s. A static-count drop is evidence that the emitted gate stream contains fewer Toffoli instructions, though the source diff is still required to identify the responsible mechanism; equality is not sufficient to establish re-sampling, because two circuit bodies can contain the same number of gates under different controls, so we classify a lower executed average as re-sampling only when the source and op-stream comparison also finds no circuit-body change apart from the identity tail. Applied to the eleven commits that a diff read alone had left *neutral*, this test reclassified eight as real gate reductions and one as verifier search; with a third padding-only submission surfaced in the wider re-audit, exactly three submissions across the full audit change no gate at all, each merely adjusting the injected DUMMY_TOFFOLIS padding rather than the circuit. Some reclassified reductions are *approximate* (a width truncation that errs on a sub-permille fraction of inputs and validates only under a co-hunted island); we count them under their gate-removing mechanism but flag the integrity concern in Section 6.

A score-weighted view. The census counts how *often* each family was the primary move; the recovered leaderboard scores (Supplemental Note C) also permit an accounting view of realized frontier movement. Let Δ_i be the ledger’s signed score change for entry i , so an improvement has $\Delta_i < 0$. The 70 largest reductions $-\Delta_i$ account for 91.6% of the gross reduction across all 400 scored entries; excluding the one that merely zeroes the injected DUMMY_TOFFOLIS gates leaves 69, whose sum is the denominator of the percentages below. This is a book-keeping decomposition of realized deltas, not a counterfactual estimate of each mechanism’s isolated effect: source commits can bundle changes, and later deltas depend on inherited earlier structure.

Under that accounting rule, the four structural families—register/scratch allocation, fusion, windowing, and uncompute strategy—receive $\approx 60\%$ of attributed frontier improvement (21%, 18%, 13%, and 8% respectively), across just 26 of the 69 included commits (under 7% of all scored commits). Register-lifecycle changes, boundary `cswap`-merge fusions, the commit landing the two-jump inversion base, and measurement-based comparator uncomputes are associated with realized deltas of hundreds of millions of score points; the two-jump landing records a delta of nearly one billion. The contrast between count and attributed delta is sharpest for fusion and windowing, which are rare by count (3.0% and 2.0% of scored commits) yet together receive nearly one third of the accounted improvement. Dead-code and redundancy elimination receives

39%, consistent with sustained late-stage tuning rather than a small number of step changes. Conversely, the families that a lever-*name* count makes look substantial shrink under a diff audit and under weighting alike: algorithmic substitution (Section B.2) is small on both axes, and the codec/encoding and constant-propagation families receive essentially no delta in this scalar-frontier accounting—their moves lie mostly on the low-width Pareto branch (Section 5.4.3), are superseded, or rarely advance the scalar incumbent.

B.1 Windowing

In the accepted-commit census (Figure 10), windowing is the primary move in 8 competition commits (2.0%), led by the wholesale landing of the two-jump inversion base and its K2 refinements; the family opens the anatomy because that small count carries the competition’s largest single frontier movement. The windowed modular shift described below is the census’s *windowed carry/shift blocks* change-type—the primary move in the remaining 5 of these 8 commits; the specific Solinas-shift commits detailed below are first documented before the competition, while the same windowed-shift primitive recurs as a competition primary move.

Windowing asks whether a batch of w consecutive loop iterations can be processed in one wider operation rather than step by step. Its closest compiler analogy is loop blocking with precomputation: a run of the body is replaced by an amortized operation controlled by a small input window. Two instances appear in the audited archive; the first changes the implementation of the circuit’s dominant cost centre.

Two-jump (K2) inversion windowing. The family’s largest instance re-blocks the circuit’s dominant cost centre, the modular inverse, so that each reversible iteration advances *two* steps of the binary-GCD recurrence rather than one; Section 5.3.1 develops the construction, its macro-step invariant, and its fixed convergence budget. Read as a taxonomy entry, it is a windowing move because the recurrence and the returned inverse are unchanged and only the iteration boundaries move, which makes it exact conditional on convergence and leaves the enclosing point-addition schedule untouched; the saving lands on average executed Toffoli. Unlike an earlier era’s naive low-bit matrix-lookup form, which a matrix-cost lower-bound test decisively closed as a moonshot, the comparator-aware two-step synthesis *landed*: it shipped as the `DIALOG_GCD` route (`src/point_add/kaliski_jump.rs`), a wholesale alternate inversion implementation that became the accepted competition base at submission `cddd5df` and was then refined by a long series of admitted submissions. A representative refinement, `DIALOG_GCD_K2_PAIR_COMPRESS` (submission `9a81017`), packs the six raw transcript bits of two consecutive K2 steps into five sidecar bits—the pair code of Section 5.3.2, licensed by a reachability constraint between one step’s shift bit and the next step’s low branch bit; it is reported validator-clean over all 9,024 shots at 1,313 logical qubits and 1,536,923 average executed Toffoli. Being an exact transcript re-encoding, it is distinct from the iteration-bound

truncations of Section B.3 and in particular from the documented approximate “396-iteration” redteam variant that this operating point deliberately excludes.

Windowed modular shift. The archive’s second windowing instance batches a no-op stretch of the Solinas reduction that normalizes products modulo $p = 2^{256} - 2^{32} - 977$. That reduction walks the set bits of $c = 2^{32} + 977$, whose nonzero positions are $\{0, 4, 6, 7, 8, 9, 32\}$, doubling and correcting at each step; between bit 9 and bit 32 the loop performs 22 consecutive `mod_double_inplace_fast` calls that reduce nothing, because no bit of c is set in that range, yet each still pays a full mod- p double-and-correct. Submission `c702290` replaces that stretch with a single classical shift-by-22 that captures the overflow into a fresh `spill` register, one sparse windowed correction that adds `spill · c` back at c ’s set-bit positions, and a single final conditional subtract; the swap-based shift itself costs no Toffoli, so the saving in average executed Toffoli is exactly the cost of the per-iteration reduction boundaries that are deleted. The substitution is exact rather than approximate and survives the evaluator’s reversibility and phase checks (Section 3.3) by a lifetime discipline: the classical correction flags are threaded through to a matching `mod_shift_right_by_k` that replays the same steps in reverse against the same qubits, so the reverse shift exactly cancels the forward one—an earlier attempt, `579db56`, failed the phase-garbage check precisely because it recomputed an independent inverse instead. The same primitive was applied at all three Solinas call sites (`c702290`, `de71b0f`, `d396bbc`), with a later refinement (`0c6235f`) exploiting the pseudo-Mersenne structure directly.

Both instances leave their algorithm in place and only re-block its schedule, which distinguishes windowing from its neighbours: fusion (Section B.6) merges exactly two adjacent operations to delete one load/unload boundary, strength reduction (Section B.7) holds the schedule fixed and swaps a single operation for a cheaper equivalent, and algorithmic substitution (Section B.2) swaps the underlying algorithm outright.

B.2 Algorithmic Substitution

Under a diff-level audit this family is the primary move in only 2 scored source commits (0.5%): a Karatsuba squarer backend swap (`28fe2f2`) and a Karatsuba x -tail-square swap (`b6d6f41`). The wired *main* modular multiplier stays schoolbook and Kaliski remains the sole inversion route for the entire competition—a name-based count over-reports this family at 61 (15.3%) because dozens of commits carry `_KARATSUBA_ / fermat_inv` lever names whose backends are never wired (default-`false` flags, `#[cfg(test)]` scaffolds), while their real dominant move is dead-code, register, or strength reduction. Both commits sit in the census’s single *multiplier/squarer swap* change-type; the schoolbook conversions and in-Kaliski reverts below predate the competition and have no primary-move instance within it.

A substitution qualifies here only when the replacement is extensionally identical over $\mathbb{F}_p$ —it must return the same residue modulo $p = 2^{256} - c$ for every admitted input, so that the enclosing point-addition schedule is unchanged and

the swap is invisible to the validator of Section 3.3 apart from its gate count—which is why these swaps primarily move average executed Toffoli rather than peak width.

Schoolbook (grid) multiplication. The baseline modular multiply is a Horner shift-and-add walk that performs $O(n)$ full conditional modular additions—a comparator, a Cuccaro add, and a Solinas reduction per bit of the multiplier. The schoolbook form instead builds the $n \times n$ array of bit-products with cheap AND gates, accumulates them into a single wide $2n$ -bit register, and reduces modulo p once at the end; the two schemes return the identical product, so the choice is a per-site cost decision, not a change to the arithmetic. The first squaring conversion (`a85389c`) saved roughly 170,000 average executed Toffoli per call, and converting the remaining `build()` multiplies to schoolbook form removed on the order of 624,000 across five commits. The wide accumulator is real ancilla, however, so this cheaper algorithm needs more live scratch and can raise the peak-width factor of the score $S = QT$ (Section 3.2) wherever two such calls are simultaneously live in one closure.

Karatsuba multiplication. Karatsuba substitutes schoolbook’s $O(n^2)$ grid for the classical three half-width-product recombination, an $O(n^{1.585})$ scheme that again returns the identical product; Section 5.3.4 gives the squaring-specialized form that survives in the final circuit, whereas the commits below are the earlier generic-multiplier swaps. Applied to all four multiply sites (`cc427db`) it lowered average executed Toffoli from 4,927,684 to 4,680,924, but its half-sum and cross-term registers add live scratch, which sharpens the same width trade-off. Pushing Karatsuba into the three multiply sites *inside* the Kaliski closure—the highest-peak-qubit region of the circuit—reduced Toffoli yet worsened the product, and `bc2f2a4` reverted exactly those sites to a scratch-lean add/subtract schoolbook variant, trading a modest Toffoli increase for a larger width reduction. The net value at any given site is therefore decided jointly with the register-allocation constraints of Section B.4 and the operation-level rewrites of Section B.7, which supplied the add/subtract schoolbook variant that won those in-Kaliski sites back.

The family’s largest instance acts on the other cost centre, the modular inverse: the whole-inverse jump/safegcd reformulation that computes the same $v^{-1} \bmod p$ through a different recurrence. In its landed form that reformulation batches two GCD steps per reversible iteration through a precomputed transition, so we classify it under Windowing (Section B.1, the 2-jump route) and record it here only as this family’s boundary case, where a change of algorithm and a change of schedule coincide.

B.3 Dead-Code and Redundancy Elimination

This is the archive’s most frequent primary move by a wide margin: 244 of the 400 scored source commits (61.0%). Comparator and register-width truncation, iteration-count and threshold tuning, and fold and carry-truncation levers

dominate, with a small dead-gate census tail. Its prevalence reflects the late-competition regime, after the circuit skeleton had stabilized and most tuning removed exact redundancies or support-conditionally truncated rarely exercised work.

Section 5.3.7 states the transformations themselves—the self-inverse cancellations, the two inductive bit-range invariants that license width truncation, and the exact and support-conditioned iteration bounds. This subsection adds what the taxonomy needs: why the family is the only genuinely dual-factor one, how its members split across two correctness models, and the commit-level record behind each.

It is the taxonomy’s genuine dual-factor family because a deleted loop iteration removes both the nonlinear work inside it *and* the registers sized to the iteration count, so a single transformation can move average executed Toffoli (Section 3.2) and peak logical width at once. Its members span two qualitatively different correctness models—*exact* deletions redundant for every admitted input, and a single *probabilistic* deletion licensed only over the validator’s fixed instance support—and that distinction, not the size of the saving, is what the ordering below foregrounds.

Mapped onto the census change-types: *self-inverse cancellation* and *negation elision* below are the *dead-gate/self-inverse removal* bucket (21); *comparator and register-width truncation* is the census’s single largest change-type (149); *iteration truncation* (exact and probabilistic) is *iteration/threshold tuning* (37); and a *fold/carry truncation* tail—fold vents and carry-tail skips, not written up separately here—accounts for the remaining 37.

Self-inverse cancellation. An involution wrapped around a helper call—a negate/act/un-negate bracket—is removed by inlining the helper’s body for the sign already on hand, so both involutions vanish from the gate stream rather than being paid for and cancelled at run time. In the Kaliski inverse plumbing the two `mod_neg_inplace_fast` calls bracketing the pair-1 modular-inverse unadd were deleted this way (0340f24, 8,584,680 → 8,583,660 average executed Toffoli), and the identical pattern on the pair-2 unadd followed one commit later (23d3fa0).

Negation elision. This reaches the same matched bracket from upstream: rather than collapsing a bracket that already exists, it chooses the operand order at a value’s source that never produces the unwanted sign—feeding Kaliski $Rx - Qx$ instead of $Qx - Rx$ and flipping the one downstream consumer’s add/subtract polarity to absorb the flip (ee1743e, 19,041,786 → 19,040,766), then propagating a single negative-native sign convention through the shared `kaliski_forward` primitive so that a negation running four times per point addition is removed once at the source (53a4aeb, 19,040,766 → 19,038,718).

Comparator and register-width truncation. The family’s single most frequent move—149 of the 244 dead-code submissions—narrows an operation to the bits that can be nonzero rather than deleting a whole loop pass. Because the two invariants of Section 5.3.7 between them cover the whole iteration range, any

operation touching bits outside the proven-nonzero prefix is an identity paid for in full; the resulting rewrites narrow an adder to its live prefix and drop a control whose two outcomes provably coincide (a `cswap` chain of $n - 1 = 255$ Toffoli-bearing Fredkin gates becoming free swaps when $v_w[0]$ is zero on both branches, `e996391`). In the source record these are the `COMPARE_BITS`, `WIDTH_MARGIN`, and `KAL_*_TRUNC_W` levers.

Iteration truncation. Where width truncation narrows an operation, iteration truncation removes whole loop passes, shrinking the gates *and* the counters and history registers indexed by the loop bound. The exact $2n - 1 = 511$ bound was established not by formal proof but by exhaustive search over the $v = 2^k$ stress class plus random and near-boundary sampling, treated as sufficient because the harness re-checks classical correctness every run—that single bound drops every iteration-indexed register by one (`c46a010`), and the larger late-iteration truncations (for example `afa876f`, $-259,080$ Toffoli) compose freely on top because each only deletes work already scoped to be a no-op; skipping a provably-unreachable Solinas correction is the same move. Every deletion in these paragraphs is exact for all admitted inputs.

Probabilistic iteration truncation. Pushing the same bound *below* its provable floor looks like the same move but changes the correctness model, since a rare input requiring more iterations would silently produce a wrong inverse. It is therefore not an all-input theorem but a validation-support-conditioned claim, licensed only over the fixed Fiat–Shamir instances the validator of Section 3.3 actually exercises and justified by two support-conditioned ingredients rather than a proof: a cited application-level tolerance argument that Shor’s algorithm absorbs on the order of one per-cent per-point-addition error [1], and the empirical zero-failure sampling reported in Section 5.3.7. Because it also shrinks the iteration-indexed history registers, `5f2adb5` is a genuine double win and among the family’s largest single reductions on either axis—yet its correctness lives entirely in how the truncation is justified: tellingly, an earlier commit trying the identical numeric bound under a *deterministic*-correctness framing was reverted, while the same change succeeded once presented as an explicit bounded-error claim.

B.4 Register and Scratch Allocation

This family is the primary move in 50 submissions (12.5%), spanning scratch reuse, ancilla-lifetime, and in-place-mutation changes. A diff-level audit finds it far more common than a lever-name count suggests because many peak-qubit moves are recorded under arithmetic-sounding lever names. In census terms, *scratch reuse* is *scratch/ancilla reuse* (17), *in-place mutation* is the like-named change-type (3), and *ancilla-lifetime windowing* is *ancilla-lifetime / venting* (30, the family’s largest bucket). *Horner evaluation* carries no census node of its own: it is a multiplier-width streaming trick whose primary move is counted elsewhere and is included here only for its mechanism.

This family is defined by changing the placement or lifetime of live storage, so its primary target is Q , the peak logical width in $S = QT$ (Section 3.2), rather than average executed Toffoli T . It is not the only family capable of moving Q : codec, fusion, and windowing changes can do so as well, and individual allocation changes can also alter T . The distinguishing mechanism is that an allocation or lifetime change can lower a maximum without removing a gate, provided it intersects every phase tied at that maximum. The paragraphs below are the local, one-register instances of that mechanism; its wholesale expression, in which complementary lifetimes let remainders, Bézout coefficients, and quotient state share two physical work registers, is the register-shared inversion of Section 5.3.5.

Scratch reuse. The mildest move: when a fresh ancilla would be allocated, borrow instead an in-scope register already provably at $|0\rangle$, or replace a hand-rolled call site with the shared primitive that already embodies the win. The correctness obligation is localized but not vacuous: the borrowed register must enter and leave at $|0\rangle$, remain unaliased while in use, and be restored before its owner returns. Handing Kaliski step 1’s `mcx3_polar` scratch to the provably-zero `b_f` (3e495b1) or reusing `add_f` at two more sites (8626f25) elides alloc/free book-keeping without altering a Toffoli, while computing a shared sub-term once into a known-zero ancilla removes a redundant compute/uncompute pair (dfa4bd0, fb60dee).

In-place mutation. This refuses the copy entirely: a routine that needs a negated, doubled, or offset version of a value it already holds walks the live register through the whole transform and restores it exactly with the inverse sequence, deleting the scratch copy, its copy-in and copy-out passes, and its free. Correctness follows from the mutation being run with `emit_inverse`-derived exact inverses, so the round trip returns the register to its entry value; aliasing the raw Kaliski inverse register into `with_kal_inv` rather than copying it removes exactly one $n = 256$ -qubit register from the peak (419a4df, 3339 $\rightarrow$ 3083), as does the companion `sum_y` deletion (93cee90, 3595 $\rightarrow$ 3339, both -256). The same discipline can even pay to *raise* the gate count: walking the Karatsuba tail through 22 in-place mod- p doublings and their inverse halvings rather than a shift-by-22 deletes the ~ 24 -qubit spill/overflow/flag lane that alone bound the peak, buying a peak-width drop at a small Toffoli cost (fe29b53)—a register move precisely because it resets Q , not T .

Horner evaluation. This attacks the multiplier’s live width by streaming: an MSB-first multiply-accumulate doubles the accumulator it is already touching rather than a separate scratch copy of the multiplicand, so the intermediate products never all need to be resident at once and the whole scratch register plus its double/halve restore walk disappears whenever the accumulator is provably zero on entry. It is correct precisely because “doubling the accumulator” only preserves meaning when that accumulator starts at $|0\rangle$, which is exactly the from-scratch multiply case, and the doublings land on `acc` that the conditional

adds already touch—no separate state to restore (**b77a311**, $\sim 130k$ Toffoli per applicable multiply; secondary strength-reduction membership).

Ancilla-lifetime windowing. This addresses the high-water mark most directly: a register whose declared scope spans a hot loop but which a peak-driving section provably never touches—because an invariant holds it at $|0\rangle$ or a known fixed value—is released across that section and reallocated fresh afterward. The release is valid when the narrowing invariant is established before the window and the register is reallocated to a fresh $|0\rangle$ (and re-flipped to its fixed value if needed) before its next use, changing no gate. Freeing **v_w** across the multiplication-heavy **body** on this argument (**1fd2656**, $3478 \rightarrow 3222$) is the single largest qubit reduction in its scan range, again an $n = 256$ -qubit drop at zero Toffoli cost.

B.5 Codec / Encoding

Classical systems also compress state; this family is separated because a reversible codec must provide a coherent decode, use, and cleanup path while trading resident quantum state for encoded metadata. It therefore spends Toffoli to buy back peak qubits. Its 7 census submissions (1.8%) are all the single *trailmix* / *K5 codec* change-type, and its landed instance in the scoring circuit is the base-5 transcript codec of Section 5.3.2, whose proof obligations—support preservation, exact decoding, and complete ancilla and phase cleanup—are exactly those that distinguish this family from classical compression.

What places the family at the taxonomy’s edge is that its most aggressive expression pushes the same trade wholesale: rather than keep the full transcript, quotient, and route values of the point-addition walk live in registers, an encoding packs them into a compact per-step description and reconstructs each value on demand, so far fewer qubits stay resident across the peak. This is the extreme end of the register-and-scratch-allocation spectrum—where scratch reuse borrows one known-zero register and windowing frees one at its true last use, the codec converts whole classes of resident arithmetic state into narrower encodings whose reversal must be replayed—and it is the mechanism behind the low-width track’s lowest-admitted-width artifact, **b6f2b0a** at 825 qubits, which correspondingly pays 489,161,900 average executed Toffoli (Table 2, Section 5.4.3).² The trade is the steep reversible time–space exchange [2,3] read through a single factor: the codec buys reductions in Q by spending in T , and the exchange rate steepens once metadata and denominator-witness lifetime, rather than ordinary scratch, become the binding constraint.

² The source record attests the codec mechanism directly—a **trailmix_ludicrous/codec.rs** module (submission **bdb1d22**) and a family of *K5* tail/head compressors that pack per-step raw state into fewer physical qubits—but the 825/489,161,900 figures are the *artifact*’s operating point (Section 5.1, Table 2), not a codec-specific measured delta; we attribute the artifact to this mechanism, not a particular qubit or Toffoli saving to the codec in isolation.

B.6 Fusion

Fusion, the primary move in 12 submissions (3.0%), merges two *adjacent* operations so that the load/unload or compute/uncompute boundary each would pay independently need never be materialized; the intervening body is left untouched, so the family moves average executed Toffoli rather than peak width. Of the paragraphs below, the first three are the census’s *gate/op fusion* change-type (7 submissions) and the fourth, *boundary cswap-merge*—collapsing two adjacent controlled swaps to delete one load/unload boundary—is the primary move in the remaining 5.

Comparator conjugation. The recurring shape is a self-inverse comparator called twice—once to compute a flag, once to uncompute it—bracketing the caller’s body, so the palindromic MAJ-only sweep runs forward and inverse twice for $4n$ Toffoli where $2n$ of carry-chain work is actually done. At **34c64b9** the step-2 “greater-than” comparator is folded into a single conjugation—one forward MAJ sweep, read the carry, run the body inline, read the carry again to clear the flag, one inverse sweep—so that $2n$ Toffoli suffice (2,096,640 fewer). The identical rewrite of the zero-test comparator at **b655093** removed a further 2,087,232.

Shared-operand scratch fusion. When two adjacent arithmetic operations share an operand, the load and unload of their common scratch register can be collapsed rather than paid twice. At **90775e2**, Kaliski step 4’s two conditional Cuccaro calls—a subtract of u from v_w then an add of r into s , both gated by the same control—were fused around one **tmp** register, the in-place transform **tmp** $\oplus=$ **add_f** $\wedge (u \oplus r)$ reshaping the loaded value into what the second operation needs instead of unloading and reloading it (524,288 fewer Toffoli).

Branch collapse. Fusion can also merge an unconditional pass with its own conditional undo. At **c413d6d**, the fast Solinas reduction’s unconditional add of the constant c followed by a conditional un-subtract—in effect both branches of an **if/else**—collapse to a single conditional add once the post-shift overflow bit is recognized as the correction control (655,864 fewer). The move is order-sensitive: the paired-addition fusion at **0bede51**, which loads the classical constant Q_x once and doubles it in place rather than loading it twice, became profitable only *after* **c413d6d** had made in-register doubling cheap.

Boundary cswap-merge. The recurring shape is a pair of width- n controlled swaps in each Kaliski iteration: step 3 orients the working registers, and step 9 commits the branch decision. Each of the two GCD channels therefore pays $2n$ Toffoli per iteration on swaps alone. The Fredkin identity $\text{cswap}(p) \cdot \text{cswap}(q) = \text{cswap}(p \oplus q)$ lets an iteration’s step-9 swap be deferred and merged with the *next* iteration’s step-3 swap: a single persistent “frame” parity qubit carries the deferred control a_k across the iteration boundary, the merged swap is driven by the combined parity $a_k \oplus a_{k+1}$ formed with a free CX, and the two width- n swaps collapse to one plus a constant-size parity fix-up. Reversibility requires

the backward pass to mirror the same deferral, so the frame qubit is cleared at the boundary that set it; because it is allocated only over that short span and is never live during the step-4 peak, the merge is peak-neutral and moves only average executed Toffoli. At `d35d6e7` the construction was wired on the (r, s) channel for roughly 274,000 fewer executed Toffoli, the largest single fusion win of the competition. Commit `112d5a4` then extended the identical mechanism to the (u, v_w) channel, restricted to a 254-iteration safe prefix. Beyond that prefix, the cheap parity correction is no longer provably valid because $u = v_w$ can occur near the walk’s terminus. The remaining submissions in this pattern only re-search that safe-iteration bound, leaving the mechanism unchanged.

Fusion is strictly a one-adjacent-pair transformation: unlike its neighbours in Section B.1, it keeps both operations and removes only the seam between exactly two of them.

B.7 Strength Reduction

A diff-level audit finds strength reduction as the *primary* move in 14 scored competition source commits (3.5%)—equivalent-but-cheaper operation swaps (negate-then-add, Solinas and NAF folds, fast-adder reroutes) that keep the algorithm’s schedule intact—so it carries a node in Figure 10. It was nonetheless chiefly a pre-competition family, and the techniques below are drawn from that earlier record, where each mechanism first becomes source-auditable. Of the census’s two change-types, *fast-adder reroute* (9) is *Fast-adder substitution* below and *solinas / NAF fold* (5) covers *Solinas reduction* and *signed-digit constant decomposition*; *negate-add substitution* and the *Litinski add-subtract* row are pre-competition techniques with no primary-move instance during the competition.

Where algorithmic substitution (Section B.2) swaps one primitive for a different one, strength reduction keeps the schedule of modular operations intact and replaces a single operation with a provably equal but cheaper reversible realization; the win is invisible to the abstract arithmetic but decisive for the gate count, and it moves average executed Toffoli rather than peak width (Section 3.2).

Solinas reduction. Because `secp256k1`’s modulus is $p = 2^{256} - c$ with $c = 2^{32} + 977$ a small, sparse 33-bit constant, any $T \in [0, 2p)$ satisfies $T \bmod p = (T \bmod 2^{256}) + c$ whenever the top-bit overflow fires (i.e. $T \geq 2^{256}$), so reduction needs only a conditional correction by the small, sparse c controlled by a single overflow bit rather than a full comparison and subtraction against the dense 256-bit p . Inlining that `add-c/branch/parity` pattern into Kaliski’s step 8 (`3536f4c`) removes 2,090,712 Toffoli. The rewrite is bit-identical because the reduction predicate is exactly the shifted-out top bit and, in the doubling case, the correction flag is uncomputed for free by parity, since the input was just doubled and p is odd; the companion single-add refinement that collapses the speculative add-then-undo pair is the branch collapse of Section B.6.

Negate-add substitution. The identity $acc - X \equiv acc + (p - X) \pmod{p}$ is free algebraically, but a conditional modular subtractor carries a comparator and underflow-correction pass that the conditional adder does not—1,536 versus 1,280 CCX per iteration in this circuit’s schoolbook multiplier—so the technique negates the subtrahend once, drives the inner loop with the cheaper add primitive, and negates back. Applied to the hottest inner loop `mod_mul_sub_qq` (`a462064`), this removes 256,016 Toffoli, a fixed one-time cost amortized against a per-iteration saving over a 256-bit multiply. The computed residue is unchanged because $p - X$ is X ’s exact modular negation and the loop reconstructs the identical accumulator.

Signed-digit constant decomposition. Multiplying a register by a fixed classical constant costs one Cuccaro add per set bit, so recoding the Solinas correction constant $977 = 0b1111010001$ (six set bits) as the four-term signed expression $977 = 2^{10} - 2^6 + 2^4 + 2^0$ collapses the contiguous run $\{6, 7, 8, 9\}$ into a single $2^{10} - 2^6$ pair and shortens the operation list without touching the algorithm or the registers (`b1d04b7`, then `a0b7de6` at the outer Solinas loop). Since a Cuccaro subtraction is merely the inverse gate sequence of an addition and costs the same, and the shift’s reverse is updated symmetrically with flipped signs, the recoded constant computes the identical product.

Litinski add-subtract schoolbook. This variant replaces each schoolbook row’s AND-compute-then-add construction with a single controlled two’s-complement add-subtract row. The former costs $2n-1$ Toffoli; the latter costs $n-1$ and needs no separate AND row, roughly halving the per-row cost. Laid inside the half-width inner multiplies of 1-level Karatsuba (`28737ac`), it is the largest single Toffoli reduction in its scan range (333,794, from 4,672,931 to 4,339,137), at a modest peak-width cost (+107 qubits) from the wider $(2n+1)$ -bit accumulator. The row still evaluates $x \cdot y$ exactly: the controlled add-subtract layering produces $2xy$ plus a fixed algebraic correction that is divided out by relabeling bits.

Fast-adder substitution. Swapping an ancilla-lean in-place Cuccaro adder for the pre-existing carry-ancilla “fast” variant shortens the gate-count-bearing path at five call sites in `mod_shift_left/right_by_k` (`384ed8a`: −17,400 Toffoli for +21 peak qubits). Because it buys Toffoli back with qubits, it is a strength reduction only when peak width has local headroom—the commit’s own check confirms the transient peak stays below the enclosing Solinas `mod_add_qq_fast` peak—so this is a boundary case whose secondary membership is in register allocation (Section B.4). The two adder variants compute the identical sum; only their scratch footprint and critical path differ.

B.8 Constant Propagation and Classical Specialization

This family is the primary move in 19 competition source commits (4.8%), all of the host-derived classical-specialization kind: commits that recognize a value as classically fixed on the validator-supported path and specialize or host it

accordingly. The two named techniques below were established before the competition, but—unlike strength reduction, which largely receded once it began—classical specialization remained a live *primary* lever throughout, which is why its competition-era count is substantial. Both techniques below are the census’s single *host-derived / classical specialization* change-type (19).

This family collects the transformations that exploit a value the circuit already knows classically—a fixed multiplier bit, a compile-time offset, or a loop branch that is provably forced on every admitted input—to specialize or simply skip reversible work; because the exploited value is genuinely fixed, the specialized circuit answers the validator of Section 3.3 on every admitted instance exactly as the general one would, and the saving lands on the average-executed-Toffoli factor (Section 3.2) at unchanged peak width.

Classical conditioning. The family’s primitive move—specializing a sub-circuit on a multiplier bit that is classical for the evaluation instance rather than held in superposition—is called *Constprop* in the source record and is developed in Section 5.3.6.

Guaranteed-prefix specialization. The family’s larger instance applies the same specialize-on-a-known-predicate idea to an entire loop prefix rather than a single gate, and in doing so it borrows from algorithmic substitution (Section B.2) as a secondary character. Kaliski’s binary almost-inverse (Section B.2) terminates when a sum invariant reaches a target value; for 256-bit `secp256k1` inputs that invariant starts above 2^{255} , so a minimum-step-count argument shows the loop *cannot* terminate during its first several iterations for any nonzero input. Those leading iterations therefore always take the same non-terminating branch, and the general branch-covering machinery a generic `kaliski_iteration` must carry is dead weight there. The optimization replaces that guaranteed-nonterminal prefix with a bespoke branch-free primitive (`kaliski_iteration_bulk_prefix3`), proven classically equivalent to the generic body restricted to the forced branch—by a companion `kaliski_equiv.rs` module checked against sampled reachable states, with a mirrored backward step—*before* it is wired into the live forward/reverse passes. Its rollout was staged: `e80edd9` first landed the specialized steps behind an opt-in flag (`KAL_BULK3_EXPERIMENT`), leaving the scored build untouched, and `91fe506` then enabled the prefix by default while extending the proven-safe span to 304 iterations, moving the scored circuit from 4,394,546 to 4,236,292 average executed Toffoli (−158,254, or 3.6%) at unchanged width; `a25978e` later restored that default on top of a 511-iteration robustness baseline and recovered the identical figure. Because the replaced iterations are proven identical to the generic body on *all* reachable nonzero inputs, this is a guaranteed-prefix argument and not the empirical, support-selected truncation of Section B.3, and it does not narrow the admitted input set.

B.9 Uncompute Strategy

Uncompute strategy is the primary move in 22 submissions (5.5%, all measurement uncompute)—a diff-level audit finds it far more common than a lever-

name count reports, because measurement-based clears (`hmr/cz_if`) are invisible in the commit subject and surface only in the diff; and the choice it governs is additionally exercised in almost every submission, because reversibility is not optional. All 22 census submissions are *measurement uncompute*—the *Measurement-based uncompute* paragraph below; the other four strategies are the compute–uncompute / Bennett discipline that reversibility forces on almost every submission but that is rarely itself the scored primary move.

Every scratch value must return to $|0\rangle$ or the run fails the reversibility check of Section 3.3, so the question this family answers is never *whether* to uncompute but only *how*; the factor its choices move is average executed Toffoli, and the units below are the menu of ways to avoid paying the reverse sweep in full.

Ancilla-clean uncompute. The baseline discipline is compute–copy–uncompute: compute an intermediate into fresh scratch, copy the wanted result out with a single `CX`, then replay the computation in reverse to drive every ancilla back to $|0\rangle$ before freeing it. This roughly doubles the gate count of the forward computation and is the $\sim 2\times$ price that replaces a banned *dirty free*: once the harness’s strict-reset check began panicking on any nonzero free (`dc8fe9d`), primitives that had “passed” by dirty-freeing a still-set ancilla had to be honestly rebuilt, as when `cmp_neq_zero` was re-expressed as a palindromic de-Morgan OR chain with an explicit reverse sweep (`266f418`). Every strategy below is an attempt to pay less than this full reverse sweep.

Compute–act–uncompute (Bennett conjugation). Rather than hand-write the reverse sweep, this strategy brackets an action between a compute block C and its exact inverse C^{-1} and makes C^{-1} mechanical: an `emit_inverse/conjugate` helper records the forward block’s gate stream, re-emits it reversed, and panics on any non-invertible op, so the teardown can never drift out of sync with the setup (`fe8875a`, which introduces the `emit_inverse/conjugate` helpers and rebuilds `cmod_half` as the exact inverse of `cmod_double`). The trade-off is that the reverse is generated for free but is only ever the recorded gates played backward—it can never substitute a cheaper equivalent reverse.

Emit-inverse. The same machinery can define an entire standalone primitive as the mechanical gate-inverse of its partner—subtract from add, halve from double—whenever the forward primitive is already clean, since a clean reversible circuit’s own gate-level inverse is correct by construction. Defining `mod_sub_qq` as `emit_inverse(mod_add_qq)` (`dad5057`) removed a hand-written negate–add–negate triplet, lowering executed Toffoli from 101,284,162 to 98,893,282 and, by deleting the `tmp` register that triplet needed to reconstruct its own operand, cutting peak width from 3,595 to 3,339 (-256 qubits)—the rare uncompute move that helps both factors of $S = QT$ (Section 3.2) at once. Its limitation is intrinsic: `emit_inverse` can only replay recorded gates backward, so the moment a cheaper non-gate-reversal inverse exists (`8212f6b`) it becomes a ceiling and the reverse must be hand-written outside the mechanism.

Measurement-based uncompute. When a scratch ancilla holds an isolated AND product that is provably unchanged at the point it must return to $|0\rangle$, Gidney’s trick replaces the second Toffoli of the textbook uncompute with a Hadamard-measure-reset (**Hmr**) and a single classically-conditioned **CZ** on the two original controls, at 0 Toffoli per bit rather than one. The cost moves off the gate-count axis onto a measurement plus a feed-forward classical bit—favourable under a validator that scores executed Toffoli—but only after the surrounding gadget is restructured so the AND product lives in its own qubit; applied to a dedicated Cuccaro carry lane it halves the adder from $2(n-1)$ to $(n-1)$ Toffoli in a single commit (**8d54b18**). Because **Hmr** is not a deterministic gate **emit_inverse** can capture and replay, this strategy is mutually exclusive with emit-inverse: reaching for it inside a machine-generated reverse forces that backward pass to be hand-written first (**8212f6b**).

Bennett pairing. The last strategy attacks not the cost of one uncompute but how many times an expensive one runs: when a value-mutating, self-restoring routine that is itself internally a Bennett pair is called twice back to back—once to transform an input in place and use it, once purely to restore it—the inner loop executes four times. Computing the transform into a fresh output register while leaving the input live, then uncomputing only that output, folds two conjugation pairs into one and roughly halves the loop executions without widening the peak, since the scratch footprint is reused sequentially rather than duplicated. Folding the Kaliski almost-inverse’s two call sites this way (**3af7f42** and the symmetric **408174a**) produced the largest single Toffoli reductions in the corpus ($98,887,162 \rightarrow 77,082,106$ and $77,082,106 \rightarrow 55,277,050$, each $-21,805,056$), and holding the Kaliski state live across the caller’s body (**8e0c9f6**) removed another wrapped forward-and-reverse pass ($31,642,202 \rightarrow 20,608,890$).

References

1. Babbush, R., Zalcman, A., Gidney, C., Broughton, M., Khattar, T., Neven, H., Bergamaschi, T., Drake, J., Boneh, D.: Securing elliptic curve cryptocurrencies against quantum vulnerabilities: Resource estimates and mitigations (2026), arXiv:2603.28846, <https://arxiv.org/abs/2603.28846>
2. Bennett, C.H.: Time/space trade-offs for reversible computation. *SIAM Journal on Computing* **18**(4), 766–776 (1989)
3. Buhrman, H., Tromp, J., Vitányi, P.: Time and space bounds for reversible simulation. *Journal of Physics A: Mathematical and General* **34**(35), 6821–6830 (2001)