

Making Submission 8e9c9a2 Compatible with Coherent Windowed Point Addition

Runtime-table ABI, quantum-circuit construction, and validation

ECDSA.Fail Supplemental Technical Note

9 August 2026

Abstract


Submission 8e9c9a2 implements mixed elliptic-curve point addition with a classically specified addend. A windowed Shor circuit instead supplies the addend coherently: a quantum address selects a row of a classical point table, the selected point is added to a quantum accumulator, and all address-dependent workspace must be removed. We modify the circuit to expose the quantum registers X, Y, J and $2 \cdot 2^w$ classical coordinate registers, condition QROM payload operations on those runtime table bits, derive $3x_J$ inside the circuit, handle the identity row coherently, and exactly replay each lookup after use. For $w = 16$, the resulting circuit uses $Q = 1,162$ logical qubits. On an independently seeded corpus shared with the original circuit, its average executed Toffoli count is $T = 1,684,160.924$ over 100,000 inputs. The any-channel error rates are 0.1910% for the windowed circuit and 0.1920% for the baseline, with no ancilla failures in either circuit and an exact paired sign-test value of $p_{\text{sign}} = 1.0$. Correlated four-call tests for every $w \in \{0, \dots, 5\}$ additionally exercise 6,144 additions and confirm cleanup after every call. The construction implements a reusable window-selected point-addition kernel on the same generic-affine domain as the source circuit; it is not a complete implementation of Shor’s algorithm.

1 Target Interface and Runtime ABI

Let $R = (x_R, y_R)$ be the quantum accumulator, let J be a coherent w -qubit address, and let $A_J = (x_J, y_J) = \mathcal{T}[J]$ be the selected point. The required map is

$$U_{\text{win}} : |J\rangle |R\rangle |0\rangle_W \mapsto |J\rangle |R + A_J\rangle |0\rangle_W, \quad (1)$$

where W contains the QROM payload, routing qubits, arithmetic scratch, and cleanup state. The address must remain coherent and unchanged, all arithmetic must depend coherently on the selected row, and no row-dependent payload, phase, or ancilla may survive the call.



The table is supplied through classical input registers. Its values are not embedded into the operation stream, while the selected row is materialized and processed coherently in qubits.

Figure 1: Runtime-table window interface implemented by the modified circuit.

The emitted register ABI is

$$(X, Y, J, \mathcal{T}_x[0], \mathcal{T}_y[0], \dots, \mathcal{T}_x[2^w - 1], \mathcal{T}_y[2^w - 1]), \quad (2)$$

where X and Y are 256-qubit coordinate registers, J is a w -qubit quantum register, and every table coordinate is a 256-bit classical input register. Thus, a $w = 16$ instance exposes $2 \cdot 2^{16}$ coordinate registers. The operation stream refers to their bits through classical conditions; replacing one shifted G - or public-point table by another changes the input data rather than the circuit topology. The supplied evaluator populates these registers with the reference table $\mathcal{T}[j] = [j]G$, but the emitted circuit is not restricted to that table.

2 Circuit-Level Modifications

2.1 Stage-local materialization

The six-stage arithmetic order of **8e9c9a2** is retained. Only the boundaries at which the addend is consumed are changed. Stages 1 and 6 jointly load the two selected coordinates through one address traversal; Stage 3 loads only x_J . Every payload is removed before the dialog-GCD peaks.




Figure 2: Stage-local lookup/use/replay schedule. The selected point is never retained through the arithmetic peak.

Stage 3 does not require a derived $3x$ table. It loads x_J into a 257-qubit temporary register, first accumulates x_J into X , reversibly doubles the temporary value modulo p , accumulates $2x_J$, reverses the doubling, and unloads the original x_J . Therefore, the update is $X \leftarrow X + 3x_J \pmod{p}$, while the ABI contains only the two source-coordinate columns. Stage 6 likewise derives reverse subtraction from the loaded x_J : for a nonidentity row it applies modular negation to X and then adds x_J , giving $x_J - X \pmod{p}$. No embedded $x_J + 1$ column is required.

2.2 Identity row and generic-affine domain

Row $J = 0$ represents the elliptic-curve identity $\mathcal{O}$. The circuit computes a coherent predicate $b = [J \neq 0]$, uses b to control the square-subtract in Stage 4 and modular negation in Stage 6, and uncomputes b at the end. When $J = 0$, the QROM payloads are zero, so Stages 1–3 leave (X, Y) at $(x_R, y_R x_R^{-1})$ after division; Stage 4 is disabled; Stage 5 reverses the division map and restores $Y = y_R$; and Stage 6 leaves $X = x_R$. Hence the complete branch returns $R + \mathcal{O} = R$ without measuring J .

The exceptional affine cases $R = A_J$ and $R = -A_J$ remain outside the circuit domain. They make the denominator $x_R - x_J$ vanish and require point doubling or an explicit identity output. This is the same generic-affine restriction used by the source circuit and Schrottenloher’s point-addition construction; it is distinct from the table’s identity row, which the modified circuit now handles.

2.3 Low-peak arithmetic

The QROM payload is live only near its consumer, but it still competes locally with carry ancillas. The modification suppresses carry venting at loaded-data sites and retains only a 64-bit upper comparison during its low-peak modular additions. This support-calibrated approximation exchanges additional Toffoli work and a small correctness risk for fewer simultaneously live carry qubits. The global maximum remains the inverse-application fold after all QROM payloads have been cleared.

3 Runtime QROM and Exact Replay

3.1 Why an XOR lookup loads a table word

For one table column $c_j \in \{0, 1\}^{256}$, QROM implements

$$L_c : |j\rangle |z\rangle \mapsto |j\rangle |z \oplus c_j\rangle. \quad (3)$$

If the payload starts in $|0\rangle$, then $0 \oplus c_j = c_j$, so $L_c |j\rangle |0\rangle = |j\rangle |c_j\rangle$. Operationally, a reversible address traversal constructs the predicate selecting row j and applies a classically conditioned CNOT to each payload position whose runtime table bit is one. By linearity,

$$\sum_j \alpha_j |j\rangle |0\rangle \mapsto \sum_j \alpha_j |j\rangle |c_j\rangle, \quad (4)$$

so the payload becomes correlated with the coherent address. Here c_j is one coordinate column, such as x_j or y_j , not the complete point A_j .

3.2 $K = 2$ QROAM and payload width

The single-column Stage 3 lookup uses two-row blocking. Write $j = 2h + \ell$, where h contains the high $w - 1$ address bits and ℓ is the least significant bit. The high-address traversal loads adjacent rows into two 256-qubit banks, and controlled swaps select the requested row:

$$|h, \ell\rangle |0\rangle_{B_0} |0\rangle_{B_1} \mapsto |h, \ell\rangle |c_{2h}\rangle_{B_0} |c_{2h+1}\rangle_{B_1} \mapsto |h, \ell\rangle |c_{2h+\ell}\rangle_{B_0} |c_{2h+1-\ell}\rangle_{B_1}. \quad (5)$$

Because $j = 2h + \ell$, $c_{2h+\ell} = c_j$ is precisely the selected word. The notation does not mean that $c_{2h+\ell}$ is the complete table point: in Stage 3 it is the x -coordinate word x_j .




Figure 3: $K = 2$ QROAM used for the Stage 3 x_J episode.

The construction does not store a 512-bit point in the 16 address qubits. Stages 1 and 6 temporarily allocate two 256-qubit payloads, and Stage 3 temporarily allocates the selected and adjacent 256-qubit banks shown in (5). The width saving is temporal: these banks are cleared before the global dialog-GCD peak instead of remaining resident throughout the point addition. Consequently, their local 512-qubit storage increases the measured peak by only 11 qubits relative to the mixed-addition circuit.

3.3 Exact replay unload

After the arithmetic consumer preserves a loaded payload, applying the same XOR lookup again gives

$$L_c^2 : |j\rangle |c_j\rangle \mapsto |j\rangle |c_j \oplus c_j\rangle = |j\rangle |0\rangle. \quad (6)$$

The reported circuit therefore reverses the $K = 2$ swaps when applicable and replays the same runtime-table traversal. This exactly removes the payload without measuring it. The traversal's unary-routing AND ancillas are removed using measurement-based uncomputation with conditional phase correction, yielding a clean coherent QROM boundary with no address-dependent residual phase. A measurement-based *payload* unload remains implemented as an optional experiment, but the final resource and correctness results use payload replay.

4 Circuit-to-Source Correspondence

These changes establish window compatibility because J is a genuine quantum register, every table-dependent payload operation is controlled by the externally supplied table bits, arithmetic consumes qubit payloads, and exact replay removes all which-row information outside J . The resulting circuit topology can therefore be reused with the shifted tables $\mathcal{T}_G^{(i)}[j] = [j2^{iw}]G$ or $\mathcal{T}_P^{(i)}[j] = [j2^{iw}]P$ required by a complete windowed scalar multiplication.

Episode	Load Toffoli	Replay Toffoli	Total	Runtime columns
Coordinate differences	65,534	65,534	131,068	x_J, y_J
Apply $3x_J$ ($K = 2$)	33,022	33,022	66,044	x_J
Output recovery	65,534	65,534	131,068	y_J, x_J
Total	164,090	164,090	328,180	six traversals

Table 1: Instrumented nonlinear QROM cost for $w = 16$. The complete stream also contains 167,772,160 classically conditioned table-payload CNOTs across loading and replay.

Circuit role	Principal source entry points	Effect
Runtime ABI	<code>trailmix_ludicrous/mod.rs</code> : <code>build_trailmix_ludicrous_ops</code>	Declares X, Y, J and the $2 \cdot 2^w$ classical coordinate registers; no table values are written into the emitted circuit.
Runtime payload controls	<code>qrom.rs</code> : <code>xor_payload</code> , <code>load_word</code> , <code>load_two_words</code> ; <code>point_add/mod.rs</code> : <code>cx_if</code>	Emits CNOT operations whose classical condition is the corresponding runtime table bit.
Stages 1 and 6	<code>ec_add.rs</code> : <code>qrom_coord_differences</code> , <code>qrom_coord_recover</code>	Loads two coordinate columns through a shared traversal, applies the coordinate maps, and immediately replays the lookup.
Runtime $3x_J$	<code>ec_add.rs</code> : <code>qrom_coord_add3x</code>	Loads only x_J , applies $x_J + 2x_J$, reverses the modular doubling, and unloads x_J .
Identity handling	<code>ec_add.rs</code> : <code>toggle_nonzero_address</code> ; <code>square.rs</code> : controlled square-subtract	Constructs $[J \neq 0]$, controls the two operations that distinguish the identity branch, and uncomputes the predicate.
Exact unload	<code>qrom.rs</code> : <code>unload_word</code> , <code>unload_two_words</code> , <code>emit_qroam2_replay_unload</code>	Reverses swaps and replays the XOR lookup to clear payload and routing qubits exactly.
Peak-width control	<code>arith.rs</code> : low-peak modular add/subtract variants	Avoids simultaneous payload and wide carry-vent allocations and uses a 64-bit upper comparison at the three QROM-use sites.
Validation only	<code>eval_circuit.rs</code> : shared-corpus generation, runtime-table initialization, paired baseline/window evaluation, durable checkpointing, and correlated sequence runner	Supplies arbitrary table values and verifies coordinates, address, incremental phase, and non-interface ancillas. These changes do not alter the emitted circuit.

Table 2: Correspondence between circuit-level changes and implementation entry points.

5 Resource Measurements

The fully emitted $w = 16$ circuit contains 213,387,745 operations. Its measured peak remains the inverse-application fold, where all QROM payloads have already been released. Table 3 compares the result with the reported 8e9c9a2 operating point. The two executed-Toffoli values are measured on their respective deterministic evaluator supports, so the table characterizes resource overhead rather than a paired error experiment.

Metric	Original 8e9c9a2	Runtime-table $w = 16$	Change
Peak logical width Q	1,151	1,162	+11 (+0.96%)
Average executed Toffoli T	1,299,453	1,684,156.770	+384,703.770 (+29.61%)
$Q \times T$	1,495,670,403	1,956,990,167	+30.84%
Static emitted operations	9,133,449	213,387,745	23.36×

Table 3: Resource overhead of coherent runtime-table selection. Table-payload CNOTs dominate the static stream and are not counted as Toffoli gates.

The final circuit uses one more qubit and more Toffoli work than the earlier fixed-table, measurement-unload prototype. That increase is the cost of closing the principal interface and cleanup gaps: table data are now runtime inputs, and every payload is removed by exact replay. The persistent $w = 16$ address itself does not imply a 16-qubit increase in Q , because the original and modified circuits have different concurrent scratch allocations at their maxima.

6 Validation

6.1 Full 16-bit evaluation

The final 16-bit circuit was evaluated on 9,024 deterministic pseudorandom cases. Each case supplies the complete runtime table $\mathcal{T}[j] = [j]G$, samples a target accumulator and coherent address, and classically computes the expected generic-affine sum. Cases with nonzero J and $x_R = x_J$ are excluded because they fall outside the declared generic-affine domain; $J = 0$ is retained.

Outcome	Count
Classical mismatches	14 / 9,024 (0.1551%)
Phase-failure shots	10 / 9,024 (0.1108%)
Ancilla-failure shots	0 / 9,024 (0%)
Shots failing any channel	20 / 9,024 (0.2216%)
Identity selections	1 / 9,024 (0.0111%)

Table 4: Validation of the final $w = 16$ runtime-table circuit. Failure categories overlap.

The identity case returned the input coordinates, preserved $J = 0$, and had no classical, phase, or ancilla failure. Exact payload replay produced no QROM cleanup failure on any shot. The residual classical and phase failures are consistent with the fixed-length dialog-GCD schedule, inherited truncated arithmetic, and the window adapter’s 64-bit upper comparison; the present experiment does not supply a coherent failure flag for these approximations.

The evaluator records one durable row per input and one commit record per 64-shot batch. On restart, it truncates any uncommitted input tail, reconstructs aggregate counts, and evaluates only unfinished batches. This checkpoint mechanism produced the complete 9,024-case ledger used above.

6.2 Paired 100,000-input random-corpus comparison

To compare the arithmetic error of the two interfaces without allowing either operation stream to select its own test set, we generated 100,000 cases from the explicit seed `ecdsafail-windowed-independent-random-100k-v1`. This seed is independent of both circuit hashes and therefore does not use the benchmark’s Fiat–Shamir sampling rule. Each case contains a random nonidentity accumulator $R = [a]G$, a uniformly sampled 16-bit address j , the addend $A_j = [j]G$, and the expected result $R + A_j$. Cases with $R = \pm A_j$ are resampled because they lie outside the common generic-affine domain, whereas $j = 0$ is retained as the identity row. Joining the two input ledgers by index found no mismatch in their address, accumulator, addend, or expected-output columns.

Metric	Original <code>8e9c9a2</code>	Runtime-table $w = 16$
Peak logical width Q	1,151	1,162
Average executed Toffoli T	1,299,440.793	1,684,160.924
$Q \times T$	1,495,656,352.743	1,956,994,993.688
Classical-output failures	154 (0.1540%)	151 (0.1510%)
Phase failures	115 (0.1150%)	117 (0.1170%)
Ancilla failures	0	0
Inputs failing any channel	192 (0.1920%)	191 (0.1910%)
Wilson 95% interval, any channel	[0.1667%, 0.2211%]	[0.1658%, 0.2200%]

Table 5: Paired evaluation on 100,000 independently seeded inputs. Failure categories overlap.

Of the 100,000 inputs, 162 fail at least one channel in both circuits, 30 fail only in the baseline, 29 fail only in the windowed circuit, and 99,779 fail in neither. A two-sided exact sign test on the 59 discordant outcomes gives $p_{\text{sign}} = 1.0$. All 151 classical-output failures of the windowed circuit are also classical failures of the baseline; the windowed circuit introduces no new classical mismatch on this corpus and removes three, including both sampled identity rows. Phase outcomes differ more because the windowed circuit introduces additional measurement-based

routing cleanup. The observed one-input difference in total failures is negligible relative to sampling uncertainty, so the experiment supports comparable empirical error rates rather than an accuracy advantage for either circuit.

On this shared corpus, coherent runtime-table selection adds 11 logical qubits (0.956%), 384,720.131 average executed Toffoli gates (29.607%), and 30.845% to the $Q \times T$ proxy. The evaluator writes one input record before committing each 64-input batch, synchronizes both files to durable storage, and binds each checkpoint to the interface, seed, requested corpus size, resource dimensions, and operation-stream fingerprint. A restarted run truncates any uncommitted input tail and skips every committed batch, allowing the experiment to resume without discarding completed cases.

6.3 Correlated multi-call tests for $w = 0, \dots, 5$

To test composition rather than isolated calls, the evaluator retains one accumulator, supplies a new coherent address, applies the same emitted kernel four times, and checks the intermediate accumulator, preserved address, phase increment, and every non-interface qubit after each call. For each $w \in \{0, \dots, 5\}$, 256 four-call sequences were evaluated, giving 1,024 additions per width and 6,144 calls in total.

w	Identity selections	Any-failure sequences	Ancilla failures	Per-call average T
0	1,024 / 1,024	0 / 256	0	1,312,343.476
1	473 / 1,024	2 / 256	0	1,321,065.429
2	253 / 1,024	0 / 256	0	1,322,636.318
3	125 / 1,024	4 / 256	0	1,324,426.755
4	63 / 1,024	0 / 256	0	1,326,497.064
5	30 / 1,024	1 / 256	0	1,328,558.680

Table 6: Correlated four-call tests. The sparse classical/phase failures reflect the inherited approximate arithmetic; no sequence leaves QROM or arithmetic ancillas live.

The $w = 0$ circuit consists entirely of identity selections and is exact on this corpus. The $w = 1$ emitter and unload path also complete normally, closing the earlier base-case crash. Most importantly, the previous failure rate proportional to 2^{-w} disappears: row zero is no longer treated as the affine point $(0, 0)$.

7 Reproduction

The implementation is maintained on branch `codex/windowed-8e9c9a2` of `jieyilong/ecdsafail-challenge`. From its worktree:

```
cargo build --release --locked --offline \
  --bin build_circuit --bin eval_circuit

WINDOWED_MODE=1 WINDOW_BITS=16 SINGLE_CCX_FANOUT_DISABLE=1 \
./target/release/build_circuit

WINDOWED_MODE=1 WINDOW_BITS=16 WINDOWED_TESTS=9024 \
WINDOWED_MAX_ERROR_RATE=1 EVAL_THREADS=8 \
EVAL_CHECKPOINT_DIR=<checkpoint-dir> \
./target/release/eval_circuit \
--note runtime-table-k16-replay-unload
```

Set `WINDOWED_SEQUENCE_CALLS=4` to select the correlated test mode; repeat emission and evaluation with `WINDOW_BITS=0, \dots, 5` to reproduce Table 6.

The paired 100,000-input experiment uses the same explicit seed and separate checkpoint directories for the two interfaces:

```
EVAL_SHARED_SEED=ecdsafail-windowed-independent-random-100k-v1 \
  MIXED_WINDOW_BITS=16 EVAL_TESTS=100000 EVAL_THREADS=8 \
  EVAL_MAX_ERROR_RATE=1 EVAL_CHECKPOINT_DIR=<baseline-dir> \
  ecdsafail run

EVAL_SHARED_SEED=ecdsafail-windowed-independent-random-100k-v1 \
  WINDOWED_MODE=1 WINDOW_BITS=16 SINGLE_CCX_FANOUT_DISABLE=1 \
  WINDOWED_TESTS=100000 WINDOWED_MAX_ERROR_RATE=1 EVAL_THREADS=8 \
  EVAL_CHECKPOINT_DIR=<windowed-dir> ecdsafail run
```

Reissuing either command validates the checkpoint manifest and resumes only unfinished batches. The compact totals accompanying this note are in `paired_100k_counts.tsv`.

8 Limitations

- **Generic-affine exceptional cases.** The identity table row is supported, but $R = A_J$ and $R = -A_J$ remain outside the circuit domain, as in the underlying generic-affine construction.
- **Approximate arithmetic.** The inherited fixed-length dialog-GCD, finite-width arithmetic, and window-specific 64-bit upper comparison produce sparse classical and phase errors. The modified circuit does not convert these events into a coherent failure flag.
- **Finite-support evidence.** The experiments establish behavior only on the stated corpora; they are not an all-input proof or a coherent full-attack error estimate.
- **Kernel rather than full Shor.** The branch implements a reusable runtime-table point-addition call. A complete attack must instantiate all shifted G and P tables, schedule all windows, and validate the complete double-scalar multiplication and Fourier/postprocessing layers.
- **Accounting convention.** Resource counts follow the challenge model. A fault-tolerant architecture may assign different costs to classical table storage, Clifford gates, measurements, and feed-forward.

9 Conclusion

The revised construction closes the principal gaps of the fixed-table prototype. It exposes the complete classical point-table ABI, conditions QROM on runtime table bits, derives temporary values within the circuit, handles the identity row, supports $w = 0$ and $w = 1$, and exactly clears every payload by replay. The resulting 16-bit kernel stays within the requested width bound at $Q = 1,162$. On the independently seeded 100,000-input corpus, it uses $T = 1,684,160.924$ average executed Toffoli gates and has a 0.1910% any-channel error rate, statistically indistinguishable from the baseline’s 0.1920% rate in the paired comparison. Its remaining limitations are finite-support arithmetic approximations, the standard generic-affine exceptional cases, and the absence of full-Shor integration rather than defects in the lookup/use/unload interface.