


Shor's ECDLP circuit with the window controls expanded

Given P and $Q = [k]P$, evaluate $f(u, v) = [u]P + [v]Q = [u + kv]P$



The window registers pass through U_f unchanged: they are coherent QROM addresses controlling additions on R , not individual CCX-style controls.

For $n = 256$ and $w = 16$ there are 32 conceptual window updates; boundary optimizations reduce this to 28 point-addition calls in the paper's `secp256k1` estimate.

Conceptual fully coherent view; the semiclassical implementation interleaves Fourier operations with the windows and recycles the control qubits.

Generic in-place mixed addition: $R' = R + A$

Seven reversible register updates implement the affine group law while recycling X and Y .

In-place addition of the selected addend A

Input: accumulator $R = (x_R, y_R)$ and selected addend $A = (x_A, y_A)$ in $E(\mathbb{F}_p)$.

Generic-case assumptions: $x_R \neq x_A$ and $x_A \neq x_{R'}$.

Output: $R' = (x_{R'}, y_{R'}) = R + A$ is returned in the same X, Y registers.

1. $X \leftarrow x_R - x_A; \quad Y \leftarrow y_R - y_A$ $\triangleright X = d_x, Y = d_y$
2. $Y \leftarrow YX^{-1}$ $\triangleright Y = \lambda$
3. $X \leftarrow X + 3x_A$ $\triangleright X = x_R + 2x_A$
4. $X \leftarrow X - Y^2$ $\triangleright X = x_A - x_{R'}$
5. $Y \leftarrow YX$ $\triangleright Y = \lambda(x_A - x_{R'})$
6. $Y \leftarrow Y - y_A$ $\triangleright Y = y_{R'}$
7. $X \leftarrow x_A - X$ $\triangleright X = x_{R'}$


Algebraic check

$$\lambda = \frac{y_R - y_A}{x_R - x_A}, \quad x_{R'} = \lambda^2 - x_R - x_A, \quad y_{R'} = \lambda(x_A - x_{R'}) - y_A.$$

After Steps 3–4, $X = (x_R - x_A) + 3x_A - \lambda^2 = x_A - x_{R'}$; Steps 5–7 then recover $y_{R'}$ and $x_{R'}$ without allocating separate output-coordinate registers.

Use inside windowed Shor

The window word selects the addend A_j ; the lookup is used directly by the mixed-addition kernel and then uncomputed.



$A_j = (x_{A_j}, y_{A_j})$ is the elliptic-curve point selected for this call. For a w -bit window,

$$A_j = T_P^j[k_j], \quad T_P^j[i] = [i 2^{jw}]P$$

for a P -scalar window, with an analogous table T_Q^j for the public-point scalar.

The complete coherent map is

$$|k_j\rangle |R\rangle |0\rangle_{\text{work}} \mapsto |k_j\rangle |R'\rangle |0\rangle_{\text{work}}, \quad R' = R + T_P^j[k_j].$$

For $n = 256$ and $w = 16$, each scalar contains 16 such windows.


A_j is not an unexplained long-lived register: its coordinates are temporary QROM work data and are returned to zero. ECDSA.Fail supplies A classically; full Shor must validate the coherent interface.

Notation is consistent across the deck: A is the selected addend, R is the input accumulator, and R'

$= R + A$ is the updated accumulator.

From the point-addition algorithm to the 8e9c9a2 pipeline

Example: $+[u_1 2^{16}]P$, with selected addend $A = [u_1 2^{16}]P$ and output $R' = R + A$




Live quantum-register state											
	DIFFERENCES		DIVISION		PREPARE		SQUARE		MULTIPLY		OUTPUT
X register	$x_R - x_A$	$\longrightarrow$	$x_R - x_A$	$\longrightarrow$	$x_R + 2x_A$	$\longrightarrow$	$x_A - x_{R'}$	$\longrightarrow$	$x_A - x_{R'}$	$\longrightarrow$	$x_{R'}$
Y register	$y_R - y_A$		λ		λ		λ		$\lambda(x_A - x_{R'})$		$y_{R'}$

Source phases in submission 8e9c9a2 (same columns as the pipeline)					
Coordinate differences t1m_coord_x_sub t1m_coord_y_sub	Dialog-GCD division t1m_inverse	Prepare X workspace t1m_coord_add3x	Specialized square t1m_square	Reverse dialog-GCD t1m_forward_multiply	Recover output t1m_coord_y_sub_final t1m_coord_rsub_final

Pipelines 2 and 5 share the dialog-GCD transcript machinery: the first produces λ , while the second replays it to multiply by X and erase the transcript.
In ECDSA.Fail, (x_A, y_A) are classical input bits; a full windowed-Shor circuit must separately implement and validate the coherent QROM interface.

Phase 1: form the coordinate differences

$$d_x = x_R - x_A \pmod{p}, \quad d_y = y_R - y_A \pmod{p}$$



Reversible-state summary

$$|x_R\rangle_X |y_R\rangle_Y |0\rangle_{C_x} |0\rangle_{C_y} \mapsto |d_x\rangle_X |d_y\rangle_Y |0\rangle_{C_x} |0\rangle_{C_y}$$


The two coordinate subtractions are independent. The classically known addend is materialized only for the duration of each arithmetic block and is then removed, so no copy of (x_A, y_A) remains live.

Source correspondence: `tlm_coord_x_sub` and `tlm_coord_y_sub` in submission `8e9c9a2`.

Phase 1 detail: optimized coordinate subtraction

The group-law step is conventional; the implementation reduces its live width and nonlinear cost.

Serialized use of one transient coordinate register



Classical-addend specialization

In submission 8e9c9a2, (x_A, y_A) is classically known. Each coordinate is copied into C with classically controlled X gates, used by one reversible subtraction, and XOR-unloaded before C is released. Thus the two 256-qubit copies are never simultaneously live.

Pseudo-Mersenne modular correction

For `secp256k1`,

$$p = 2^{256} - F, \quad F = 2^{32} + 977.$$

The overflow or borrow correction can therefore be folded through the low 53 bits rather than by a second full-width reduction. The final conditional correction compares a selected 19-bit most-significant window.

Vented carry cleanup

The adder stores selected carries in temporary ancillas during the forward pass, then removes them by X -basis measurement and conditional phase feedback:

$$|c\rangle \xrightarrow{H, M_X} m, \quad \text{CZ}^m \text{ restores the phase.}$$


Each vent exchanges one temporary qubit for one fewer carry-uncompute Toffoli, which is favorable in this off-peak phase.

Interpretation. Phase 1 does not introduce a new point-addition identity. Its savings come from source-level liveness control, arithmetic specialized to p , and measurement-assisted uncomputation. The classical loading shown here is valid for the challenge's fixed-addend interface; a coherent windowed-Shor implementation instead requires a separately validated QROM lookup/use/unlookup construction.

Source correspondence: the sequential `coord_addsub` calls and `mod_sub_vented` implementation in submission 8e9c9a2.

Phase 2: compute the slope by dialog-GCD division

$\lambda = d_y d_x^{-1} \pmod{p}$ while restoring $X = d_x$ and clearing the transcript




$$(u, v; X_B, Y_B) : (p, d_x; 0, d_y) \xrightarrow{\text{walk and inverse-transpose replay}} (1, 0; \lambda, 0), \quad \lambda = d_y d_x^{-1} \pmod{p}$$

Each Jump-2 step records whether subtraction, swapping, and a second halving occurred. Only five of the eight nominal three-bit symbols are reachable, which permits the compressed dialog representation.

Source correspondence: `mod_mul_inverse_in_place(..., Direction::Inverse)`, `forward_gcd_jump`, and `reverse_gcd_jump`.

Phase 2 detail: the Jump-2 Euclidean step

Up to two binary-GCD halvings are represented by one reversible control symbol



Only five symbols are reachable

$s_2 = 0 \Rightarrow b = 1$, because the value remaining after the first halving is odd;
 $b = 0 \Rightarrow s = 0$, because swapping is enabled only on a subtraction step.

(1, 0, 0)

(1, 1, 0)

(0, 0, 1)

(1, 0, 1)

(1, 1, 1)

Why the recorded symbol is sufficient

The walk applies a data-dependent linear update M to (u, v) . The paired arithmetic applies the inverse transpose M^{-T} to (X_B, Y_B) , preserving the division relation. The bit s_2 selects whether this paired update contains one or two powers of two, while b and s select subtraction and operand order.




$$\text{walk: } (u, v) \xrightarrow{M} (u', v') \iff \text{paired update: } (X_B, Y_B) \xrightarrow{M^{-T}} (X'_B, Y'_B)$$

Compared with a one-step binary-GCD schedule, Jump-2 packages the mandatory halving and a possible second halving into one recorded iteration. Submission [8e9c9a2](#) executes 261 such iterations and replays their symbols in reverse order to recover the quotient while restoring the Euclidean workspace.

Implementation correspondence: `forward_gcd_jump`, `apply_step_reverse`, `apply_step_forward`, and `reverse_gcd_jump`.

Phase 2 detail: K5 codec and streamed dialog replay

Compress the five-way Jump-2 control alphabet, then decode only the symbols needed by the active replay step



Phase 2 detail: width scheduling and qubit reuse

The GCD state becomes smaller as the dialog grows, and short-lived workspaces are fitted into the released lanes

Scheduled active widths

$$w_i = \text{SCHED_J2}[i]$$

At iteration i , the implementation retains only the active w_i -bit slices of u and v . High-order lanes whose values are known to be zero are released, so shifts, swaps, and conditional subtraction become progressively narrower.

shrinking (u, v) creates
room for the growing tape τ

Truncated comparison windows

$$k_i = \text{cmp_window}(i, w_i) < w_i$$

The swap decision compares only a scheduled portion of the active operands. This reduces comparator work and scratch, but the selected windows are calibrated under the benchmark’s validation model rather than justified here as an all-input exact replacement.

smaller comparison window
 $\Rightarrow$ fewer gates and ancillas

Temporary hosting and cleanup

The known bits $u_0 = 1$ and $v_0 = 0$ can be parked while dormant, exposing their physical lanes to temporary state. Already-live registers also serve as dirty “vent” space for chunked adders, with every borrowed value restored before its owner becomes active again.

Measurement-assisted uncomputation clears selected AND and carry ancillas using classical phase feedback, shortening their ownership intervals without changing the logical map.

Combined live-width model

$$Q_i \approx \underbrace{512}_{\text{Bezout pair}} + \underbrace{|\tau_i|}_{\text{compressed dialog}} + \underbrace{2w_i}_{u_i, v_i} + \underbrace{q_{\text{work}, i}}_{\text{comparison, carry, and control}}$$


The central scheduling objective is to make the decrease in $2w_i$ absorb the growth of $|\tau_i|$, while hosting and measurement-based cleanup keep $q_{\text{work}, i}$ below the same peak.

Exact reversible mechanisms include support-set encoding, streamed replay, known-bit parking, restoration of borrowed dirty lanes, and phase-corrected measurement uncomputation. The chosen narrowing and comparison schedules carry the finite-support evidential status of the submitted benchmark circuit.

Implementation correspondence: SCHED_J2, cmp_window, park_known_one, park_known_zero, with_dirty_vent_pool, and clear_and.

Phase 3: prepare the X workspace

$$X = d_x \mapsto d_x + 3x_A = x_R + 2x_A \pmod{p}, \qquad Y = \lambda$$



Why the transformation prepares the square step


$$d_x + 3x_A = (x_R - x_A) + 3x_A = x_R + 2x_A$$

The following phase subtracts λ^2 , producing $x_R + 2x_A - \lambda^2 = x_A - x_{R'}$. The submitted route uses its configured low-cost modular-adder variant; the classical multiplier does not occupy quantum storage outside this phase.

Source correspondence: `tlm_coord_add3x`, `classical_times3_mod_q`, and `coord_add3x`.

Phase 4: subtract the specialized modular square

$$X \mapsto X - \lambda^2 = x_A - x_{R'} \pmod{p}, \quad p = 2^{256} - F, \quad F = 2^{32} + 977$$



Identity implemented by the streamed square

$$\lambda^2 = A + 2^{128}(C - A - B) + 2^{256}B \equiv A + 2^{128}(C - A - B) + FB \pmod{p}$$


The signed updates shown above sum to $-\lambda^2$. Each of A , B , and C is built, accumulated into X , and reversibly unbuilt before the next product is retained; $Y = \lambda$ is never overwritten.

The reduction of the $2^{256}B$ term to FB exploits the pseudo-Mersenne form of the secp256k1 prime and avoids a generic full-width modular-square workspace.

Source correspondence: *mod_square_sub_pm_secp256k1_symmetric*, including the C , A , and B build/apply/unbuild phases.

Phase 5: multiply by replaying the dialog-GCD circuit in reverse

$$Y = \lambda \mapsto \lambda X = \lambda(x_A - x_{R'}) \pmod{p} \text{ while preserving } X$$



Why the division machinery is reused


Multiplication by a quantum value must remain reversible. The source therefore runs the same in-place division construction in the opposite direction: the GCD path for X is recorded, replayed onto $(0, \lambda)$, and then erased.

$$|X\rangle_X |\lambda\rangle_Y |0\rangle_Z \mapsto |X\rangle_X |\lambda X\rangle_Y |0\rangle_Z$$

Source correspondence: `mod_mul_inverse_in_place(..., Direction::Forward)` and reverse replay with `apply_step_forward`.

Phase 6: recover the affine output coordinates

$$y_{R'} = \lambda(x_A - x_{R'}) - y_A, \quad x_{R'} = x_A - (x_A - x_{R'})$$



Final reversible map

$$|x_A - x_{R'}\rangle_X |\lambda(x_A - x_{R'})\rangle_Y \mapsto |x_{R'}\rangle_X |y_{R'}\rangle_Y$$

The fused reverse subtraction uses $\overline{X} + (x_A + 1) = x_A - X \pmod{2^{256}}$ and then applies the secp256k1 modular correction.
All loaded classical coordinates, carry state, and temporary reduction data are removed before returning the two quantum output registers.

Source correspondence: `t1m_coord_y_sub_final`, `t1m_coord_rsub_final`, and the enabled `fused_mod_rsub_vented_loaded` route.